Supplementary Material: From Feed to Video: Coordinating Multi-Scenario Recommendation with Stackelberg Games

Wenjing Li*, Jiang Rong†, Zhenzhe Zheng*, Fan Wu*
*Shanghai Jiao Tong University
Shanghai, China
†Xiaohongshu
Beijing, China
{wenjingli, zhengzhenzhe}@sjtu.edu.cn, fwu@cs.sjtu.edu.cn
rongjiang@xiaohongshu.com

***Abstract*—This supplement documents an implementation snapshot associated with the Feed–Video recommendation study. It provides source-code origins, configuration inventories, generic input descriptions, environment information, and detailed reproducibility gaps. The model interfaces and methodological implementation details are presented in the main paper. The package contains code and documentation only; it contains no datasets or model checkpoints. The archived scripts have unresolved execution and evaluation differences that are described explicitly. Consequently, the package should be read as an inspectable implementation snapshot, not as a verified reproduction of the manuscript's industrial results. The main paper remains a separate 12-page document followed by its references, with its theoretical proofs retained in the main text.**

I. Artifact Scope and Organization

The package includes twelve Python files adapted from the recovered implementation in `code/legacy/`. They cover the three Stackelberg optimizer variants, joint and hierarchical baselines, the causal-inference baseline, model definitions, JEV+ evaluation, confidence-interval utilities, ablations, and ablation plotting. Input-column identifiers and input filenames have been replaced with generic semantic names for public distribution. The computational logic and numeric settings are retained; the files are not byte-identical copies of the internal sources.

The historical project README contains the same Feed CTR improvements for SGD, Mirror Descent, and Bi-level optimization as the manuscript (20.00%, 7.59%, and 5.06%), and the same corresponding JEV+ improvements (19.74%, 17.90%, and 23.94%). This establishes a connection to the study, but does not identify the precise source revision, data split, or random seeds that produced each table. No industrial experiment was rerun while assembling this supplement.

The shared `models.py` was recovered from the parent project directory because the smaller code bundle imports it without including it. The other parent-directory additions are the causal-inference, ablation, and plotting scripts. Copying these files does not resolve every missing dependency; the remaining issues are listed in Section IV.

TABLE I
Included implementation artifacts. All filenames are relative to the legacy-code directory.

File	Role in the recovered project
`SGD.py`	Independent and Stackelberg SGD comparison
`Mirror_descent.py`	Adam-based comparison
`bi_level.py`	RMSprop-based comparison
`Base_MARL.py`	Joint Feed/Video training baseline
`Hierarchical_RL.py`	Hierarchical baseline
`Causal_Inference.py`	Causal-inference baseline
`models.py`	Shared model and dataset classes
`jev_computation.py`	Journey construction and JEV+ utilities
`jev_stackelberg_comparison.py`	JEV+ comparisons across variants
`confidence_intervals.py`	Bootstrap utilities
`ablation_study.py`	Incremental model-component analysis
`exp_plots.py`	Ablation plot from embedded aggregate values

No raw, sampled, anonymized, or synthetic dataset files are distributed. The package also excludes private logs, checkpoints, pilot output directories, and later rebuttal-result directories. The public code uses generic field names rather than internal schema identifiers. This supplement makes no commitment to release industrial data.

II. Configuration Inventory and Input Interface

The main paper's Section VIII, "Implementation and Reproducibility," now contains the shared-model architecture, historical-feature inputs, training and ranking defaults, input-processing conventions, JEV+ implementation behavior, and scope of artifact verification. The following inventory records script-level settings. Public input names describe user and item identifiers, engagement signals, Feed dwell time, and Video watch time; they do not disclose the internal schema. They

TABLE II
ARCHIVED TRAINING DEFAULTS (EPOCHS / LEARNING RATE / BATCH SIZE) AND EVALUATION LIMITS; NOT CERTIFIED FINAL-RUN CONFIGURATIONS.

Script (optimizer)	Defaults	Test users
SGD (SGD)	5 / 0.01 / 64	10,000-ID range
Mirror Descent (Adam)	5 / 0.01 / 64	10,000-ID range
Bi-level (RMSprop)	100 / 0.01 / 64	10,000-ID range
Joint baseline (Adam)	100 / 0.001 / 64	Up to 10,000
Hierarchical (Adam)	100 / 0.001 / 256	10
Causal (Adam)	3 / 0.001 / 256	10
Ablation (SGD)	5 / 0.01 / 64	Script-specific
JEV utilities (varied)	3 / 0.01 / 64	Caller-dependent

describe the archived snapshot rather than a certified final experimental configuration.

The Bi-level epoch count is imported from the joint-baseline module. The JEV utilities use SGD, Adam, or RMSprop depending on the selected variant. Most scripts expect `feed_interactions.csv` and `video_interactions.csv` in the working directory; the hierarchical and causal-inference scripts instead use paths relative to its parent. No input files are distributed. The main paper explains the split, chronology, journey-linkage, and metric differences that affect the interpretation of this interface.

III. Environment and Experiment Entry Points

The local environment inspected during packaging uses Python 3.9.6, PyTorch 2.7.1, NumPy 2.0.2, pandas 2.3.1, scikit-learn 1.6.1, Matplotlib 3.9.4, and SciPy 1.13.1. The included requirements file records these versions. It is not a recovered lockfile for the original experiment runs, and it does not capture production hardware or serving infrastructure.

The declared dependencies can be installed from the supplement directory with:

```
python -m pip install -r requirements.txt
```

The experiment entry points and their roles are listed in Table I. Running them requires the private input files, the working-directory assumptions described above, and resolution of the missing interfaces in Section IV. Installing the dependencies alone does not make the archived experiments executable. Several scripts load data and begin training at import time.

IV. Known Reproducibility Gaps

A. Execution and Dependency Gaps

The JEV comparison imports `models.train_model`, which is absent from the recovered shared model file. Mirror Descent and Bi-level scripts import `generate_confidence_report` and `compare_metrics_with_ci`, which are absent from the included confidence-interval utility. Many scripts perform data loading and training at top level. Restoring a missing function by guessing an implementation could change the experiment, so the computational logic has been retained and the missing interfaces documented.

B. Evaluation Semantics

The bootstrap utility samples user IDs with replacement and then stores them as dictionary keys. Repeated selections are collapsed, so sample multiplicity is not preserved. Some ranking-metric implementations construct the ideal NDCG list by sorting the retrieved top-k relevance values rather than deriving it from the full relevant candidate population. Together with the split and JEV differences described in the main paper, these observations prevent certification of the paper's confidence intervals or benchmark values from the recovered snapshot alone.

C. Model Branches and Experiment Provenance

When historical features are absent, the shared Video model constructs a new linear layer inside its forward method. Supplying historical features avoids that branch, but does not establish correctness of other model paths. The archived settings also differ in test-user limits, training epochs, scoring bonuses, and feature use. A confirmed final experiment revision is needed to determine which of these paths produced the submitted results.

No final-table run manifest, complete seed list, saved model state, timing environment, or online A/B assignment and logging artifact is included. The ablation plotting script contains aggregate percentages in its source. Recreating its plot would reproduce a visualization of those values, not the experiment that generated them. Existing pilot and rebuttal outputs were not substituted for the main-paper results.

V. Relationship to the Main Submission

Section VIII of the main paper now contains the implementation and verification explanations. This supplement retains inventories, environment information, and detailed limitations. It adds no empirical claims or appendix to the main document. Complete reproduction still requires the final experiment revision, missing helpers, and configuration, with any corrections accompanied by validation. All distributed files remain free of dataset records.